

simple_reinforcement_learning_exploration.R

vaishnavimyadam

2020-06-25

```
## Vaishu Myadam (vmyadam1208@gmail.com)
## June, 2020

## Exploring the reinforcement learning package with grid world and tic tac toe (Q-learning)
## Using official documentation

# Necessary libraries:

library(ReinforcementLearning)

# Tic tac toe:

data("tictactoe")
head(tictactoe, 10)

##      State Action NextState Reward
## 1  ....X.B      c7  ....X.B      0
## 2  ....X.B      c6  ...B.XX.B      0
## 3  ...B.XX.B     c2  .XBB.XX.B      0
## 4  .XBB.XX.B     c8  .XBBBXXB      0
## 5  .XBBBXXB     c1  XXBBBXXB      0
## 6  ....X.B      c1  X...B....      0
## 7  X...B....     c4  X..XB.B..      0
## 8  X..XB.B..     c3  XBXXB.B..      0
## 9  XBXXB.B..     c8  XBXXBBBX.      0
## 10 XBXXBBBX.     c9  XBXXBBBX      0

control <- list(alpha = 0.2, gamma = 0.4, epsilon = 0.1)
model <- ReinforcementLearning(tictactoe, s = "State", a = "Action", r = "Reward", s_new = "NextState",
results <- computePolicy(model)
head(results)

## .XBB..XB XXBB.B.X. .XBB..BXX BXX...B.. ..XB..... XBXXB...
##      "c1"      "c5"      "c5"      "c4"      "c5"      "c9"

summary(model)

##
## Model details
## Learning rule:      experienceReplay
## Learning iterations: 1
## Number of states:   1893
## Number of actions:  9
## Total Reward:      5449
```

```

##
## Reward details (per iteration)
## Min: 5449
## Max: 5449
## Average: 5449
## Median: 5449
## Standard deviation: NA

# Grid world:

states <- c("s1", "s2", "s3", "s4")
actions <- c("up", "down", "left", "right")
envie <- gridworldEnvironment
print(envie)

## function (state, action)
## {
##   next_state <- state
##   if (state == state("s1") && action == "down")
##     next_state <- state("s2")
##   if (state == state("s2") && action == "up")
##     next_state <- state("s1")
##   if (state == state("s2") && action == "right")
##     next_state <- state("s3")
##   if (state == state("s3") && action == "left")
##     next_state <- state("s2")
##   if (state == state("s3") && action == "up")
##     next_state <- state("s4")
##   if (next_state == state("s4") && state != state("s4")) {
##     reward <- 10
##   }
##   else {
##     reward <- -1
##   }
##   out <- list(NextState = next_state, Reward = reward)
##   return(out)
## }
## <bytecode: 0x7fd970f750a8>
## <environment: namespace:ReinforcementLearning>

data <- sampleExperience(N = 1000, env = envie, states = states, actions = actions)
head(data)

##   State Action Reward NextState
## 1    s3  down    -1         s3
## 2    s1  down    -1         s2
## 3    s4 right    -1         s4
## 4    s3  left    -1         s2
## 5    s4  down    -1         s4
## 6    s1   up     -1         s1

control <- list(alpha = 0.1, gamma = 0.1, epsilon = 0.1)
model <- ReinforcementLearning(data, s = "State", a = "Action", r = "Reward", s_new = "NextState", control = control)
computePolicy(model)

##      s1      s2      s3      s4

```

```
## "down" "right"      "up" "right"
```

```
summary(model)
```

```
##
## Model details
## Learning rule:          experienceReplay
## Learning iterations:    1
## Number of states:       4
## Number of actions:      4
## Total Reward:           -450
##
## Reward details (per iteration)
## Min:                    -450
## Max:                    -450
## Average:                -450
## Median:                 -450
## Standard deviation:     NA
```

```
# Prediction
```

```
data_to_be_predicted <- data.frame(State = c("s1", "s2", "s1"), stringsAsFactors = FALSE)
data_to_be_predicted$OptimalAction <- predict(model, data_to_be_predicted$State)
data_to_be_predicted
```

```
##      State OptimalAction
## 1      s1           down
## 2      s2           right
## 3      s1           down
```

```
# See what happens when you add new data
```

```
new_data <- sampleExperience(N = 1000,
                             env = envie,
                             states = states,
                             actions = actions,
                             actionSelection = "epsilon-greedy",
                             model = model,
                             control = control)
modified_model <- ReinforcementLearning(new_data,
                                         s = "State",
                                         a = "Action",
                                         r = "Reward",
                                         s_new = "NextState",
                                         control = control,
                                         model = model)
plot(modified_model)
```

Reinforcement learning curve

